

Chapter 4 Notes

- I. The **ps** command shows only the processes that belong to the current user and that are running on the current terminal.

A. \$ ps

PID	TTY	TIME	CMD
3081	pts/0	00:00:00	bash
3209	pts/0	00:00:00	ps

- B. The GNU ps command that's used in Linux systems supports three different types of command line parameters:

1. Unix-style parameters, which are preceded by a dash

a) Examples:

(1) -A/-e: Shows all processes

(2) -f: Displays a full format listing

(a) Information columns:

- (i) UID: The user responsible for launching the process
- (ii) PID: The process ID of the process
- (iii) PPID: The PID of the parent process (if a process is started by another process)
- (iv) C: Processor utilization over the lifetime of the process
- (v) STIME: The system time when the process started
- (vi) TTY: The terminal device from which the process was launched
- (vii) TIME: The cumulative CPU time required to run the process
- (viii) CMD: The name of the program that was started

(3) -l: Displays a long listing

2. BSD-style parameters, which are not preceded by a dash

3. GNU long parameters, which are preceded by a double dash

- II. The **top** command displays process information similarly to the ps command, but it does it in real-time mode

A. Information columns:

1. PID: The process ID of the process
2. USER: The user name of the owner of the process
3. PR: The priority of the process
4. NI: The nice value of the process
5. VIRT: The total amount of virtual memory used by the process
6. RES: The amount of physical memory the process is using
7. SHR: The amount of memory the process is sharing with other processes
8. S: The process status (D = interruptible sleep, R = running, S = sleeping, T = traced or stopped, or Z = zombie)
9. %CPU: The share of CPU time that the process is using
10. %MEM: The share of available physical memory the process is using
11. TIME+: The total CPU time the process has used since starting
12. COMMAND: The command line name of the process (program started)

B. By default, when you start top, it sorts the processes based on the %CPU value

C. Each interactive command is a single character that you can press while top is running and changes the behavior of the program

III. Linux process signals:

- A. HUP: Hangs up
- B. INT: Interrupts
- C. QUIT: Stops running
- D. KILL: Unconditionally terminates
- E. SEGV: Produces segment violation
- F. TERM: Terminates if possible
- G. STOP: Stops unconditionally, but doesn't terminate
- H. TSTP: Stops or pauses, but continues to run in background
- I. CONT: Resumes execution after STOP or TSTP

IV. The **kill** command allows you to send signals to processes based on their process ID (PID)

- A. By default, the kill command sends a TERM signal to all the PIDs listed on the command line
- B. To send a process signal, you must either be the owner of the process or be logged in as the root user

- C. The -s parameter allows you to specify other signals (either using their name or signal number)
 - 1. # kill -s HUP 3940
- V. The **killall** command is a powerful way to stop processes by using their names rather than the PID numbers
 - A. Allows you to use wildcard characters
- VI. Before you can use a new media disk on your system, you must place it in the virtual directory. This task is called **mounting**
 - A. Most Linux distributions have the ability to automatically mount specific types of removable media
 - 1. A **removable media device** is a medium that can be easily removed from the PC, such as CD-ROMs and USB memory sticks
 - B. The command used to mount media is called **mount**
 - 1. mount -t type device directory
 - 2. By default, the mount command displays a list of media devices currently mounted on the system
 - 3. The mount command provides four pieces of information:
 - a) The device filename of the media
 - b) The mount point in the virtual directory where the media is mounted
 - c) The filesystem type
 - d) The access status of the mounted media
 - 4. To manually mount a media device in the virtual directory, you must be logged in as the root user or use the sudo command to run the command as the root user.
 - 5. The -o option allows you to mount the filesystem with a comma-separated list of additional options
 - a) ro: Mounts as read-only
 - b) rw: Mounts as read-write
 - c) user: Allows an ordinary user to mount the filesystem
 - d) check=none: Mounts the filesystem without performing an integrity check
 - e) loop: Mounts a file

- C. To remove a removable media device, you should never just remove it from the system. Instead, you should always **unmount** it first
 - 1. The command used to unmount devices is **umount**
 - a) `umount [directory | device]`
 - 2. If any program has a file open on a device, the system won't let you unmount it
- VII. The **df** command allows you to easily see what's happening on all the mounted disks
 - A. The command displays the following:
 - 1. The device location of the device
 - 2. How many 1024-byte blocks of data it can hold
 - 3. How many 1024-byte blocks are used
 - 4. How many 1024-byte blocks are available
 - 5. The amount of used space as a percentage
 - 6. The mount point where the device is mounted
 - B. The `-h` parameter shows the disk space in human-readable form, usually as an M for megabytes or a G for gigabytes
- VIII. The **du** command shows the disk usage for a specific directory
 - A. By default, the command displays all the files, directories, and subdirectories under the current directory and it shows how many disk blocks each file or directory takes
 - B. Parameters:
 - 1. `-c`: Produces a grand total of all the files listed
 - 2. `-h`: Prints sizes in human-readable form, using K for kilobyte, M for megabyte, and G for gigabyte
 - 3. `-s`: Summarizes each argument
- IX. The **sort** command sorts data
 - A. By default, the command sorts the data lines in a text file using standard sorting rules for the language you specify as the default for the session
 - B. By default, the sort command interprets numbers as characters and performs a standard character sort, producing output that might not be what you want
 - 1. To solve this problem, use the `-n` parameter, which tells the sort command to recognize numbers as numbers instead of characters and to sort them based on their numerical values

- C. If you use the -M parameter, the sort command recognizes the three-character month nomenclature and sorts appropriately
 - D. Use the -t parameter to specify the field separator character, and use the -k parameter to specify which field to sort on
 - E. The -r option also sorts the values in descending order, so you can easily see what files are taking up the most space in your directory
- X. The **grep** command searches either the input or the file you specify for lines that contain characters that match the specified pattern
 - A. `grep [options] pattern [file]`
 - B. The output from grep is the lines that contain the matching pattern
 - C. Parameters:
 - 1. -v: Reverse the search (output lines that don't match the pattern)
 - 2. -n: Find the line numbers where the matching patterns are found
 - 3. -c: Count of how many lines contain the matching pattern
 - 4. If you need to specify more than one matching pattern, use the -e parameter to specify each individual pattern
 - D. The **egrep** command is an offshoot of grep, which allows you to specify POSIX extended regular expressions, which contain more characters for specifying the matching pattern
 - E. The **fgrep** command is another version that allows you to specify matching patterns as a list of fixed-string values, separated by newline characters
- XI. The **zip** utility allows you to easily compress large files (both text and executable) into smaller files that take up less space
 - A. The **gzip** utility is the most popular compression tool used in Linux
 - 1. Package includes:
 - a) gzip for compressing files
 - (1) You can specify more than one filename or even use wildcard characters to compress multiple files at once
 - b) gzcat for displaying the contents of compressed text files
 - c) gunzip for uncompressing files
- XII. By far the most popular archiving tool used in Unix and Linux is the **tar** command
 - A. `tar function [options] object1 object2 ...`
 - B. Was originally used to write files to a tape device for archiving
 - C. Functions:

1. -A: Appends an existing tar archive file to another existing tar archive file
2. -c: Creates a new tar archive file
3. -d: Checks the differences between a tar archive file and the filesystem
4. -r: Appends files to the end of an existing tar archive file
5. -t: Lists the contents of an existing tar archive file
6. -u: Appends files to an existing tar archive file that are newer than a file with the same name in the existing archive
7. -x: Extracts files from an existing archive file

D. Options:

1. -C dir: Changes to the specified directory
2. -f file: Outputs results to file (or device) file
3. -j: Redirects output to the bzip2 command for compression
4. -p: Preserves all file permissions
5. -v: Lists files as they are processed
6. -z: Redirects the output to the gzip command for compression